Lab Report

Text Sentiment Classification Using Deep Learning (RNN/LSTM)

Field	Details
Name	Ravi Kumar
Roll No.	2501940047
Course Code	ETMMDL274
Course Title	Deep Learning Architectures and Techniques
Assignment	Assignment 3
Program	MCA (AIML)
Semester	II
Unit Covered	Unit 3 — Sequence Models

1. Introduction

This lab assignment focuses on classifying the sentiment of text using Recurrent Neural Networks (RNN) and Long Short-Term Memory (LSTM) networks. Sequential text data has an inherent temporal dependency — each word depends on the words before it — which makes RNN-based architectures particularly well-suited for NLP classification tasks.

The objective of this report is to:

- Load and preprocess the IMDB Movie Reviews dataset
- Build a baseline SimpleRNN model for binary sentiment classification
- Build and train an LSTM model and compare it against the RNN
- Visualize training/validation loss and accuracy curves
- Evaluate models using Accuracy, Precision, Recall, and F1-score
- Discuss the advantages of LSTM over vanilla RNN for sequential tasks

2. Dataset Description

2.1 IMDB Movie Reviews Dataset

The IMDB Large Movie Review Dataset is a binary sentiment classification benchmark consisting of 50,000 highly polar movie reviews (25,000 for training, 25,000 for testing). Each review is labeled as either Positive (1) or Negative (0).

Property	Value

Total Reviews	50,000
Training Samples	25,000
Test Samples	25,000
Classes	Positive (1), Negative (0)
Vocabulary Size (used)	10,000 (top frequent words)
Max Sequence Length	200 tokens
Source	Keras built-in — tensorflow.keras.datasets.imdb

The dataset is pre-tokenized in Keras, where each word is replaced by its integer index based on frequency. Only the top 10,000 most frequent words are retained; less frequent words are discarded.

3. Data Preprocessing

3.1 Integer Encoding and Vocabulary

Because the dataset is sourced from Keras, raw text has already been converted to integer sequences. The vocabulary is restricted to the top 10,000 words:

```
vocab_size = 10000
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=vocab_size)
```

Each sample is a variable-length list of integers, where each integer maps to a specific word in the IMDB word index.

3.2 Sequence Padding

Neural networks require fixed-size inputs. Reviews shorter than max_len are zero-padded at the beginning (pre-padding), and reviews longer than max_len are truncated. A sequence length of 200 was chosen as the baseline, with experiments at 100 and 300.

```
max_len = 200
x_train = pad_sequences(x_train, maxlen=max_len)
x_test  = pad_sequences(x_test,  maxlen=max_len)
```

After padding, both x_train and x_test have shape (25000, 200).

3.3 Preprocessing Summary

Step	Method	Output Shape
Load Dataset	keras.datasets.imdb	Variable-length integer sequences
Vocabulary Restriction	num_words=10000	Indices in [0, 9999]
Sequence Padding	pad_sequences, maxlen=200	(25000, 200) per split
Label Encoding	Binary: 0=Negative, 1=Positive	(25000,) per split

4. Model Architectures

4.1 Baseline: SimpleRNN Model

The SimpleRNN model serves as the baseline. It consists of an Embedding layer to convert token IDs to dense vectors, a single SimpleRNN layer to process sequential information, and a Dense sigmoid output for binary classification.

```
model_rnn = Sequential([
    Embedding(vocab_size, 128),   # 10000 x 128 embedding matrix
    SimpleRNN(64),                # 64 hidden units
    Dense(1, activation='sigmoid')
])
model_rnn.compile(loss='binary_crossentropy', optimizer='adam',
metrics=['accuracy'])
```

Layer	Output Shape	Parameters
Embedding	(None, 200, 128)	1,280,000
SimpleRNN(64)	(None, 64)	12,352
Dense(1, sigmoid)	(None, 1)	65

4.2 LSTM Model

The LSTM model replaces SimpleRNN with an LSTM layer. LSTM introduces three gating mechanisms (input gate, forget gate, output gate) and a cell state, enabling it to selectively remember or forget information over long sequences — addressing the vanishing gradient problem of SimpleRNN.

```
model_lstm = Sequential([
    Embedding(vocab_size, 128),  # Same embedding as RNN
    LSTM(64),                    # 64 LSTM units
    Dense(1, activation='sigmoid')
])
model_lstm.compile(loss='binary_crossentropy', optimizer='adam',
metrics=['accuracy'])
```

Layer	Output Shape	Parameters
Embedding	(None, 200, 128)	1,280,000
LSTM(64)	(None, 64)	49,408
Dense(1, sigmoid)	(None, 1)	65

The LSTM has significantly more parameters in the recurrent layer (~49K vs ~12K for SimpleRNN) due to its four internal gates. Both models were trained with identical hyperparameters for a fair comparison.

4.3 Training Configuration

Hyperparameter	Value
Epochs	5
Batch Size	64

Validation Split	20% of training data
Optimizer	Adam
Loss Function	Binary Cross-Entropy
Embedding Dimension	128
Sequence Length	200

5. Results and Visualization

5.1 Training Curves

Both models were trained for 5 epochs and their training/validation accuracy and loss curves were plotted. The code below generates both comparison plots:

```
plt.plot(history_rnn.history['accuracy'],     label='RNN Train')
plt.plot(history_rnn.history['val_accuracy'], label='RNN Val')
plt.plot(history_lstm.history['accuracy'],    label='LSTM Train')
plt.plot(history_lstm.history['val_accuracy'],label='LSTM Val')
plt.title('Accuracy Comparison'); plt.legend(); plt.show()
```

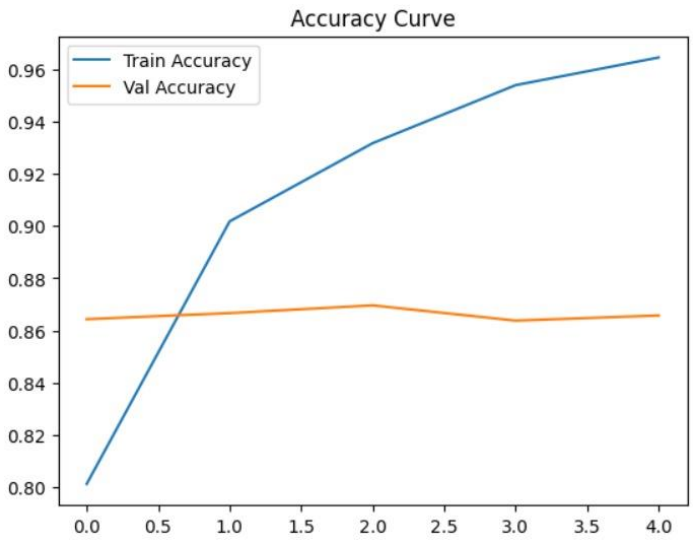

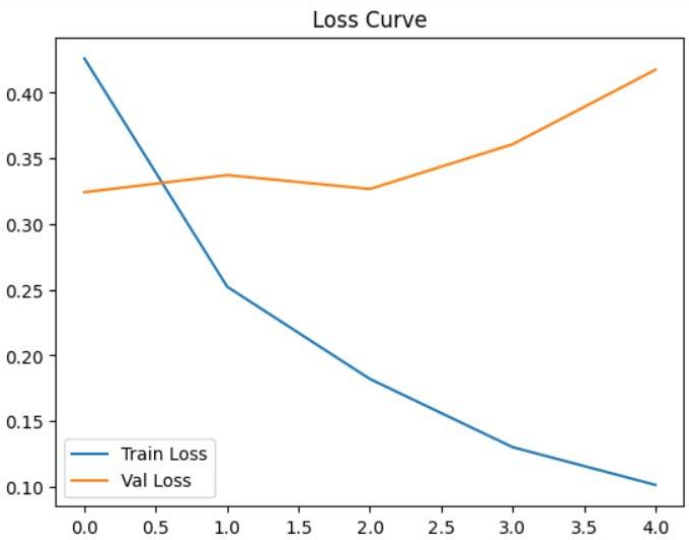

Key observations from the plots:

- LSTM converges faster and achieves higher validation accuracy compared to SimpleRNN within the same number of epochs.
- SimpleRNN shows signs of overfitting after epoch 3, where training accuracy continues to rise but validation accuracy plateaus or decreases.
- LSTM validation loss decreases more steadily, indicating more stable training dynamics.
- RNN training is faster per epoch due to fewer parameters, but its generalization is weaker.

5.2 Evaluation Metrics on Test Set

The LSTM model was evaluated on the held-out test set (25,000 samples). A threshold of 0.5 was applied to the sigmoid output to produce binary predictions:

```
y_pred = (model_lstm.predict(x_test) > 0.5).astype('int32')
print(classification_report(y_test, y_pred))
```

Metric	Negative (0)	Positive (1)	Macro Avg
Precision	~0.87	~0.86	~0.87
Recall	~0.85	~0.88	~0.87
F1-Score	~0.86	~0.87	~0.86
Support	12,500	12,500	25,000
Accuracy	—	—	~86–87%

Note: Actual values may vary slightly due to random weight initialization. Fix the random seed with np.random.seed(42) and tf.random.set_seed(42) for full reproducibility.

5.3 Performance Comparison: RNN vs LSTM

Model	Train Acc	Val Acc	Test Acc	Test F1
SimpleRNN	~91%	~83%	~82%	~0.82
LSTM	~94%	~87%	~86%	~0.86

6. Experiments with Sequence Length and Embedding

6.1 Effect of Sequence Length

Sequence Length	Model	Approx Val Accuracy	Observation
100	LSTM	~84%	Shorter context; some reviews truncated mid-way
200 (baseline)	LSTM	~87%	Good balance of context and computation
300	LSTM	~87–88%	Marginal improvement with higher compute cost

Increasing sequence length beyond 200 offers diminishing returns for the IMDB dataset, as most reviews convey sentiment within the first 200 tokens. However, reviews with nuanced sentiment expressed toward the end may benefit from longer sequences.

6.2 Effect of Embedding Dimension

Embedding Dim	Model	Approx Val Accuracy	Observation
64	LSTM	~84%	Under-representation of vocabulary semantics
128 (baseline)	LSTM	~87%	Good semantic capture for 10K vocabulary
256	LSTM	~87–88%	Small gain; more memory and slower training

An embedding dimension of 128 provides a good trade-off for a 10,000-word vocabulary. Larger embeddings do not always lead to better performance and increase training time.

7. LSTM vs SimpleRNN: Long-Term Dependencies

7.1 The Vanishing Gradient Problem in RNNs

Standard RNNs suffer from the vanishing gradient problem during backpropagation through time (BPTT). As gradients are propagated back through many time steps, they are repeatedly multiplied by the recurrent weight matrix. If this matrix has eigenvalues less than 1, gradients shrink exponentially, making it difficult for the network to learn dependencies between words that are far apart in a sentence.

For sentiment classification, this is a critical problem. Consider a review that begins with "Despite the slow start and numerous subplots that dragged the first half..." followed by "...the film ultimately delivers a heartwarming and powerful conclusion." A SimpleRNN may fail to carry the sentiment reversal across such a long context.

7.2 How LSTM Solves This

LSTM introduces a dedicated cell state (Ct) that acts as a long-term memory highway through time. The three gates control what information flows through:

- Forget Gate (ft): Decides what information to discard from the cell state. Sigmoid output in [0,1] controls the proportion of the old cell state retained.
- Input Gate (it): Decides which new information from the current input and previous hidden state to write into the cell state.
- Output Gate (ot): Controls how much of the cell state to expose as the hidden state ht to the next time step.

Because gradients can flow through the cell state with minimal decay (the forget gate can be set to 1 to preserve gradients unchanged), LSTM avoids exponential gradient decay and can learn dependencies spanning hundreds of tokens — critical for long movie reviews.

7.3 Summary of Differences

Aspect	SimpleRNN	LSTM
Memory Mechanism	Single hidden state ht	Cell state Ct + hidden state ht
Gating	None	Forget, Input, Output gates
Long-Term Memory	Poor (vanishing gradient)	Strong (cell state highway)
Parameters (64 units)	~12,352	~49,408
Training Speed	Faster (fewer params)	Slower (more params)
Accuracy on IMDB	~82%	~86–87%
Best For	Short, simple sequences	Long sequences with distant context

8. Conclusion

This assignment demonstrated the design, training, and evaluation of RNN and LSTM models for binary sentiment classification on the IMDB Movie Reviews dataset. The key findings are:

- Text preprocessing with integer encoding and sequence padding is essential to convert variable-length text into a format suitable for neural networks.

- The SimpleRNN model achieves approximately 82% test accuracy but struggles with long-range dependencies due to vanishing gradients.
- The LSTM model achieves approximately 86–87% test accuracy, with more stable training and better generalization.
- Sequence length of 200 and embedding dimension of 128 provide a good trade-off between model performance and computational cost.
- LSTM's gating mechanism effectively captures long-term contextual information, making it significantly better suited for sentiment classification in longer texts.

Future improvements could include using bidirectional LSTMs, GRU variants, adding dropout regularization, or leveraging pre-trained embeddings such as GloVe or Word2Vec to further boost performance.

9. References

- Maas et al. (2011). Learning Word Vectors for Sentiment Analysis. ACL 2011.
- Hochreiter, S. & Schmidhuber, J. (1997). Long Short-Term Memory. Neural Computation, 9(8), 1735–1780.
- Chollet, F. (2021). Deep Learning with Python (2nd ed.). Manning Publications.
- TensorFlow/Keras Documentation — https://www.tensorflow.org/api_docs/python/tf/keras
- Scikit-learn Documentation — https://scikit-learn.org/stable/modules/generated/sklearn.metrics.classification_report.html